

ABSTRACT

Addressing the significant threat that plant diseases pose to global food security and the limitations of traditional, manual detection methods, this project develops an automated diagnostic system using a Convolutional Neural Network (CNN), implemented with the TensorFlow framework. By training on a comprehensive dataset of plant leaf images, the CNN model learns to automatically extract and analyze distinguishing visual features, such as spots, texture, and color anomalies, to accurately identify and classify a range of plant diseases. The primary objective is to create a robust and efficient tool that provides a rapid and reliable diagnosis from a simple leaf image. This technology is designed to serve as an accessible solution for early disease detection, enabling timely intervention to mitigate crop loss, optimize treatment, and enhance overall productivity with AI suggestions.

CASE STUDY

PROJECT NAME: Plant Disease Detection and AI-Based Suggestion System.

INTRODUCTION:

The health of a crop is central to agricultural success, yet diseases—often manifesting first on the leaves—pose a constant and significant threat. Current detection methods are manual, time-consuming, and often leading to costly misdiagnoses and delayed treatment. By the time a disease is identified from visible symptoms on the leaves, it may have already spread, causing substantial crop losses. This project aims to solve this problem by developing an AI-powered system that can instantly identify plant diseases from a simple photo of a leaf. By providing specific, actionable suggestions for treatment, our system will empower users to proactively manage crop health, reduce losses, and increase yields for a more sustainable future.

Objectives:

- 1. CNN-Based Disease Identification:** To develop and train a robust Convolutional Neural Network (CNN) model capable of accurately identifying and classifying various plant diseases from leaf images.
- 2. Actionable Advisory System:** To build an AI-powered advisory module that provides practical recommendations for treating and preventing identified diseases.
- 3. User-Friendly Interface:** To design and implement an intuitive and user-friendly application (e.g., mobile or web) that allows for effortless image submission and instant results.
- 4. Data Collection and Knowledge Base Expansion:** To establish a system for continuous data collection and feedback from users, thereby improving the model and advisory content over time.
- 5. Sustainable Agricultural Impact:** To contribute to more sustainable and efficient farming by providing timely disease insights that help reduce crop loss and minimize the use of chemical treatments

IMPLEMENTATION

Data Preparation & Curation:

- This stage involves obtaining and organizing the dataset we'll use to train our plant disease prediction model.
- **Download the dataset from Kaggle:** We use the Kaggle to download the "plantvillage-dataset". This process confirms successful download and extraction.
- **Extract the dataset:** The downloaded file is a zip archive, which we extract to access the image files.
- **Explore the dataset structure:** We examine the extracted folders to understand how the data is organized. The dataset is divided into three main directories: 'color', 'segmented', and 'grayscale'. We'll use the 'color' images. We confirm the number of classes is 38 and explore examples of class directories.

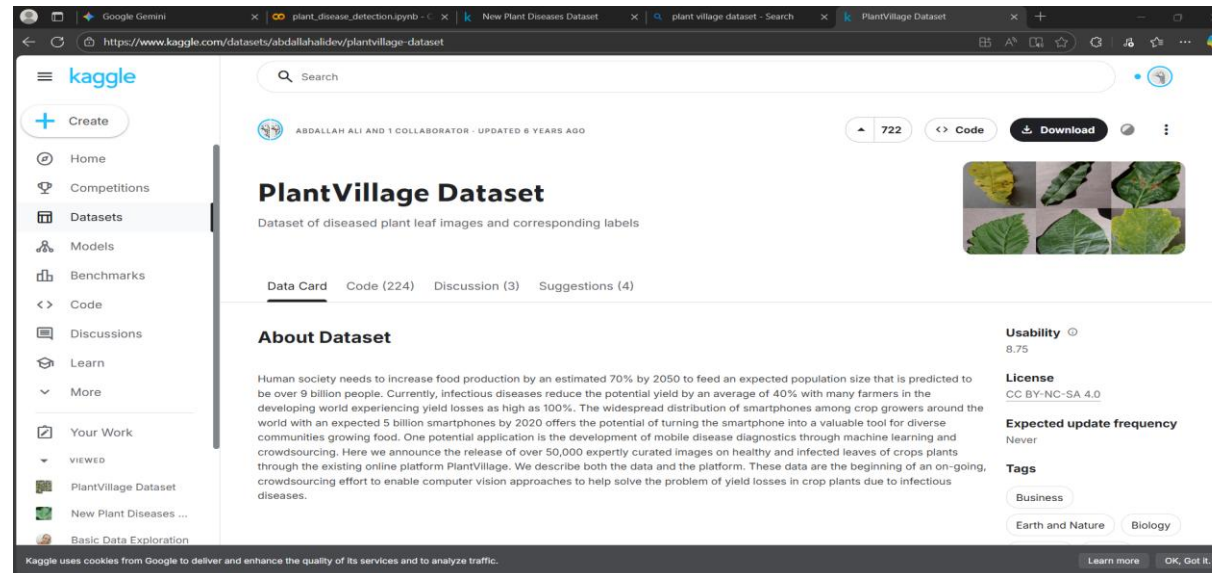


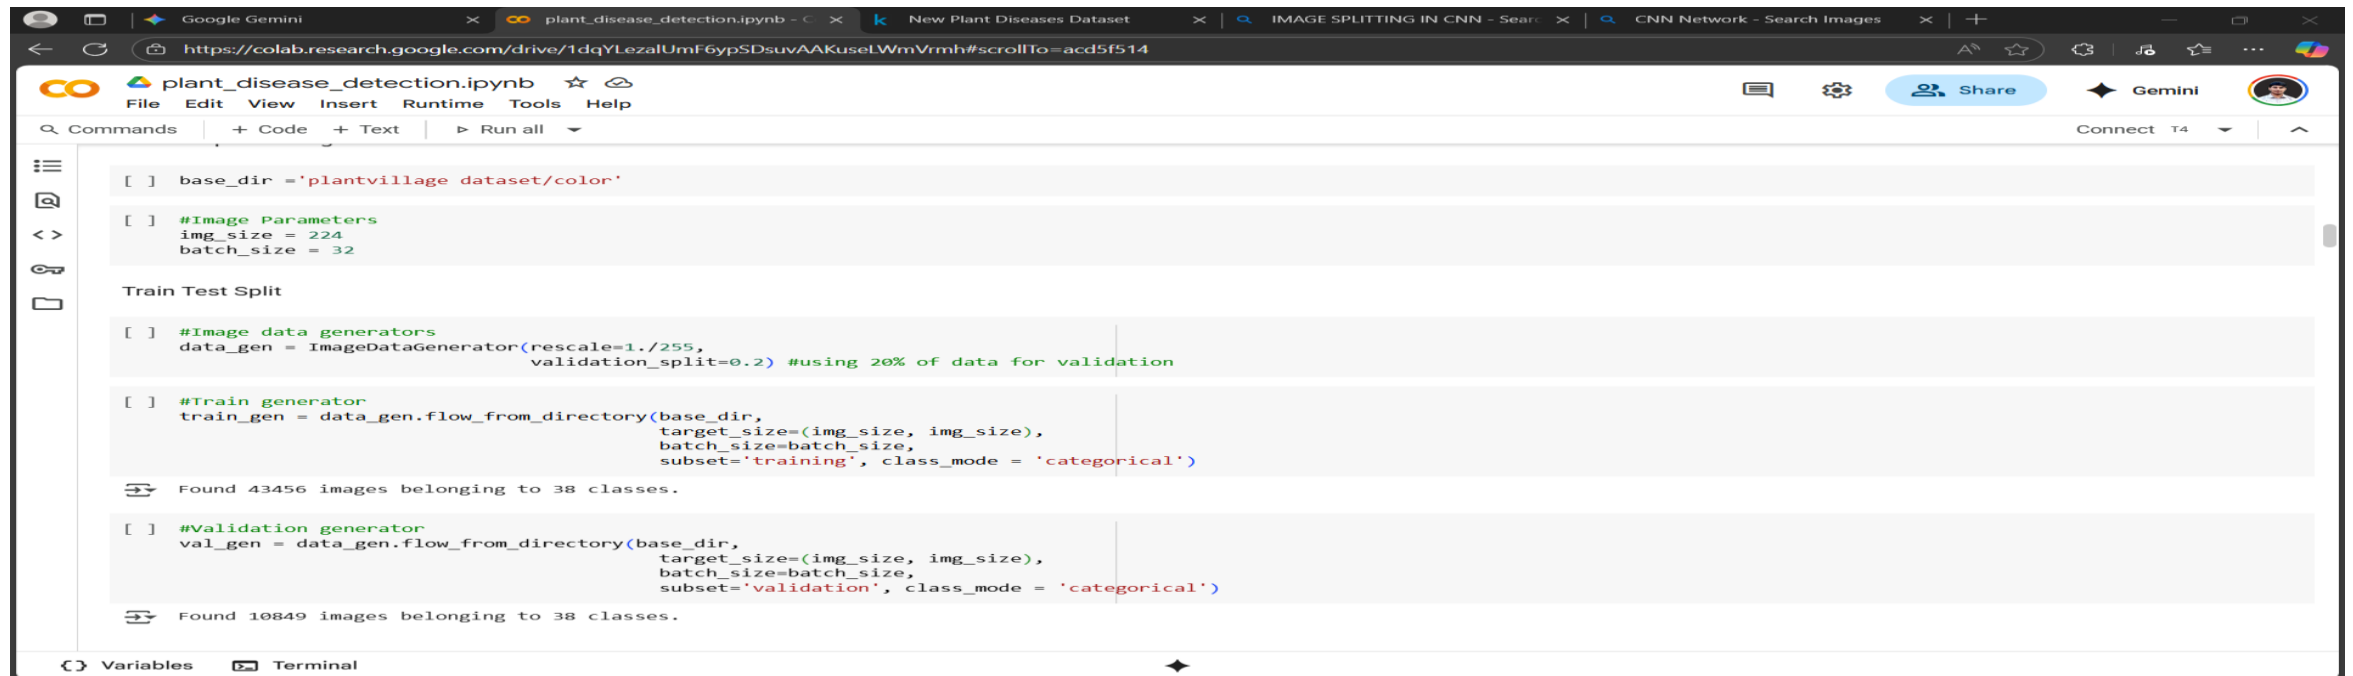
Image Preprocessing & Splitting

In this step, we prepare the images for input into our Convolutional Neural Network (CNN) and divide the dataset into training and validation sets.

Define Image Parameters: We set the target size for all images (224x224 pixels) and the batch size for training (32 images per batch).

Create Image Data Generators: We use *ImageDataGenerator* from *Keras* to handle the image data, rescaling pixel values to [0, 1] and splitting the data into 80% training and 20% validation sets.

Generate Training and Validation Data Batches: The *flow_from_directory* method reads images, resizes them, applies rescaling, and organizes them into batches. This confirms we found 43,456 images for training and 10,849 images for validation, and provides the *train_gen* and *val_gen* objects for training and evaluation.



```
[ ] base_dir = 'plantvillage dataset/color'

[ ] #Image Parameters
img_size = 224
batch_size = 32

Train Test Split

[ ] #Image data generators
data_gen = ImageDataGenerator(rescale=1./255,
                              validation_split=0.2) #using 20% of data for validation

[ ] #Train generator
train_gen = data_gen.flow_from_directory(base_dir,
                                         target_size=(img_size, img_size),
                                         batch_size=batch_size,
                                         subset='training', class_mode = 'categorical')

Found 43456 images belonging to 38 classes.

[ ] #Validation generator
val_gen = data_gen.flow_from_directory(base_dir,
                                       target_size=(img_size, img_size),
                                       batch_size=batch_size,
                                       subset='validation', class_mode = 'categorical')

Found 10849 images belonging to 38 classes.
```

Model Architecture Definition

Here, we design the structure of our Convolutional Neural Network (CNN) for image recognition.

Initialize Sequential Model: We create a Sequential model, a linear stack of layers.

Add Convolutional Layers: We add Conv2D layers (with 32, 64, and 128 filters) to detect image features.


Add Pooling Layers: MaxPooling2D layers reduce spatial dimensions and retain important information.

Add Dropout Layers: Dropout layers are included to prevent overfitting.

Flatten Layer: The Flatten layer converts 2D feature maps to a 1D vector. (i.e. 1D Array)

Add Dense Layers: Dense layers learn to classify features.

Output Layer: The final Dense layer has 38 neurons with ***softmax*** activation for class probability distribution. The ***model.summary()*** provides a summary of the defined architecture, including layers, output shapes, and parameter counts.

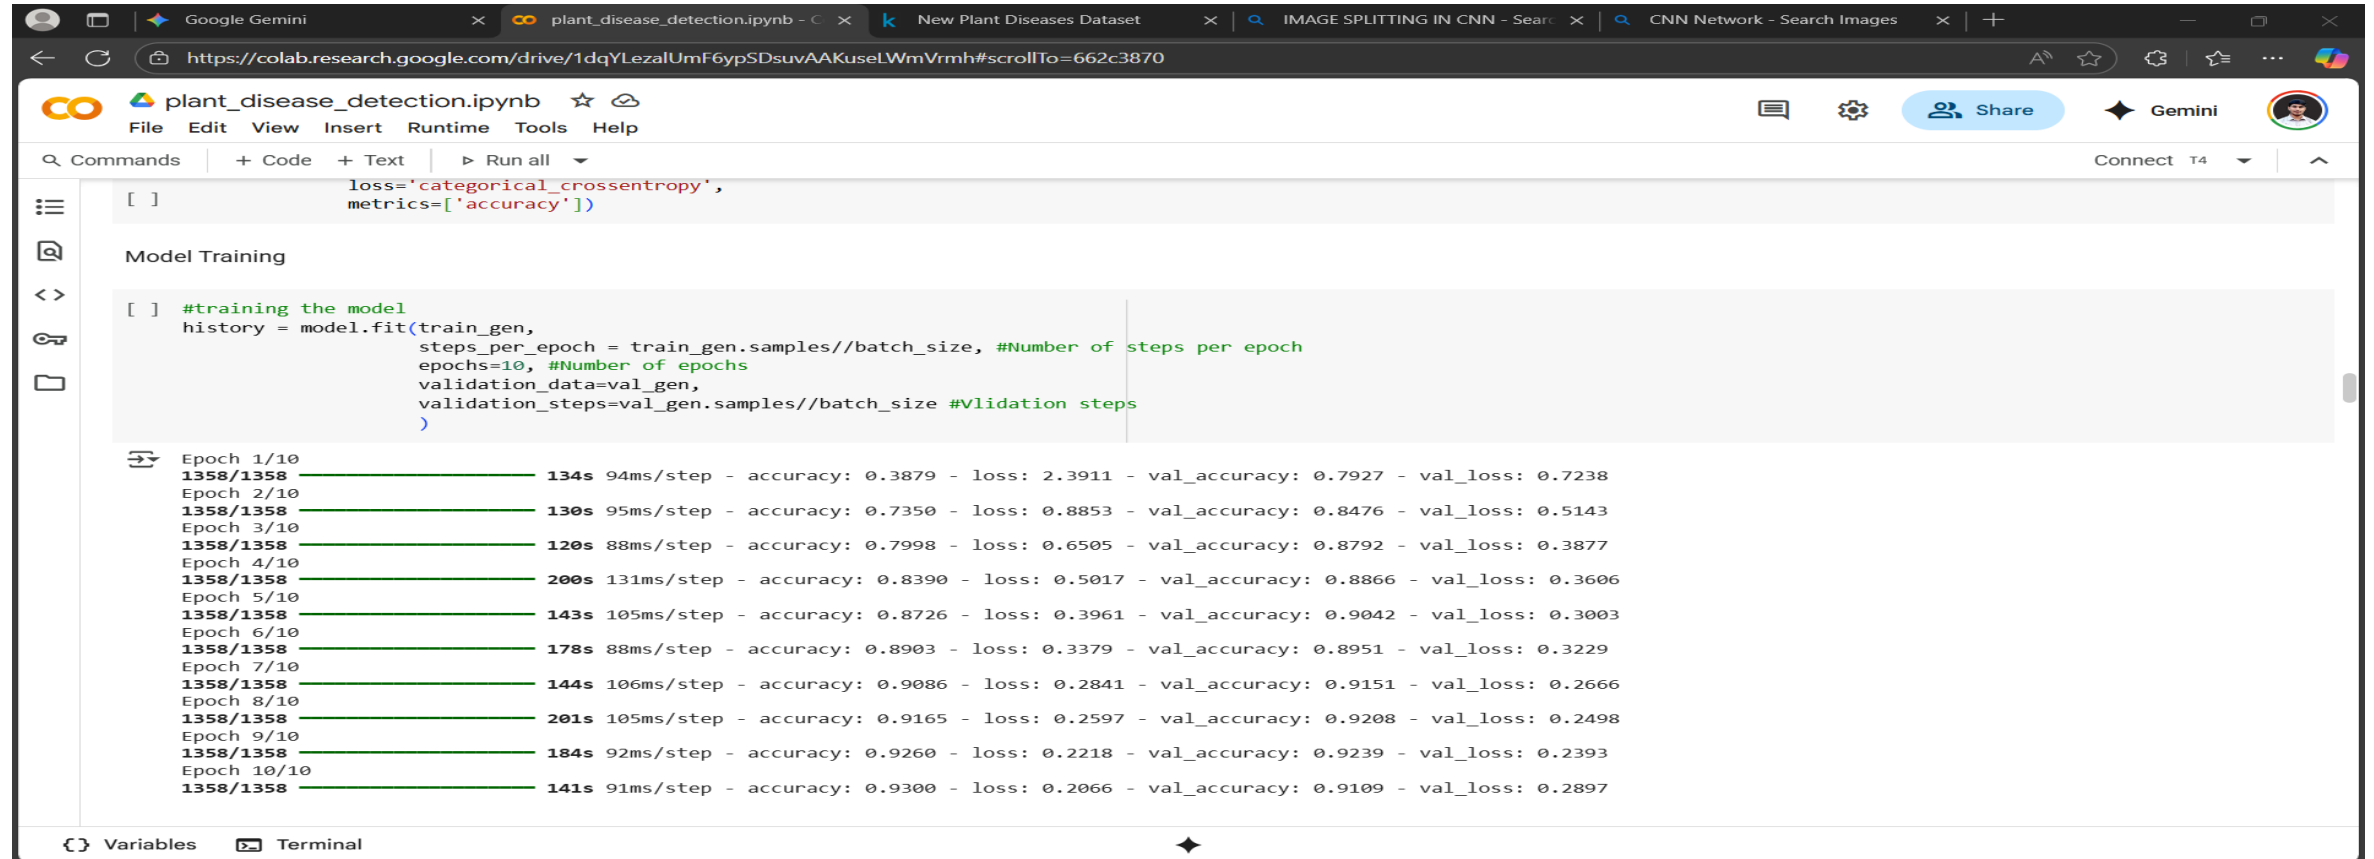


Model Training

In this phase, the CNN learns to identify plant diseases.

Compile the Model: We compile the model with the **adam optimizer**, categorical_crossentropy loss function, and accuracy metric.

Train the Model: We train the model using **model.fit()** with train_gen and val_gen, specifying steps_per_epoch, epochs, validation_data, and validation_steps. This process outputs training progress information for each epoch, including loss and accuracy, and provides a history object containing these metrics.



```
[ ] loss='categorical_crossentropy',
    metrics=['accuracy'])
```

Model Training

```
[ ] #training the model
history = model.fit(train_gen,
                    steps_per_epoch = train_gen.samples//batch_size, #Number of steps per epoch
                    epochs=10, #Number of epochs
                    validation_data=val_gen,
                    validation_steps=val_gen.samples//batch_size #Vlidity steps
                    )
```

Epoch	1/10	2/10	3/10	4/10	5/10	6/10	7/10	8/10	9/10	10/10
1358/1358	134s	130s	120s	200s	143s	178s	144s	201s	184s	141s
94ms/step	95ms/step	88ms/step	131ms/step	105ms/step	88ms/step	106ms/step	105ms/step	92ms/step	91ms/step	
accuracy: 0.3879	accuracy: 0.7350	accuracy: 0.7998	accuracy: 0.8390	accuracy: 0.8726	accuracy: 0.8903	accuracy: 0.9086	accuracy: 0.9165	accuracy: 0.9260	accuracy: 0.9300	
loss: 2.3911	loss: 0.8853	loss: 0.6505	loss: 0.5017	loss: 0.3961	loss: 0.3379	loss: 0.2841	loss: 0.2597	loss: 0.2218	loss: 0.2066	
val_accuracy: 0.7927	val_accuracy: 0.8476	val_accuracy: 0.8792	val_accuracy: 0.8866	val_accuracy: 0.9042	val_accuracy: 0.8951	val_accuracy: 0.9151	val_accuracy: 0.9208	val_accuracy: 0.9239	val_accuracy: 0.9109	
val_loss: 0.7238	val_loss: 0.5143	val_loss: 0.3877	val_loss: 0.3606	val_loss: 0.3003	val_loss: 0.3229	val_loss: 0.2666	val_loss: 0.2498	val_loss: 0.2393	val_loss: 0.2897	

Variables Terminal

Model Evaluation

After training, we evaluate the model's performance on the validation dataset.

Evaluate the Model: We use ***model.evaluate()*** with val_gen to calculate the loss and accuracy on the validation set, providing the validation loss and accuracy scores (we got 91% validation accuracy here) .

Visualize Performance: We plot the training and validation accuracy and loss over epochs using the history object, visualizing the model's learning progress.




```
[ ] val_loss, val_acc = model.evaluate(val_gen, steps=val_gen.samples//batch_size)
    print("Validation Loss:", val_loss)
    print("Validation Accuracy:", val_acc)
```




```
→ 339/339 ————— 23s 68ms/step - accuracy: 0.9146 - loss: 0.2819
Validation Loss: 0.2897346615791321
Validation Accuracy: 0.9108591675758362
```


1. Plot representing the Accuracy of Model at each Epoch



2. Plot representing the Loss of Model at each Epoch



Predictive System & Saving

In this final stage, we create a system for making predictions and save the model.

Create Prediction Function: We define `predict_image_class` to load, preprocess, and predict the class of a new image using the trained model and class indices mapping.

Create Class Indices Mapping: We create a `class_indices` dictionary from `train_gen.class_indices` that maps numerical indices to class names.

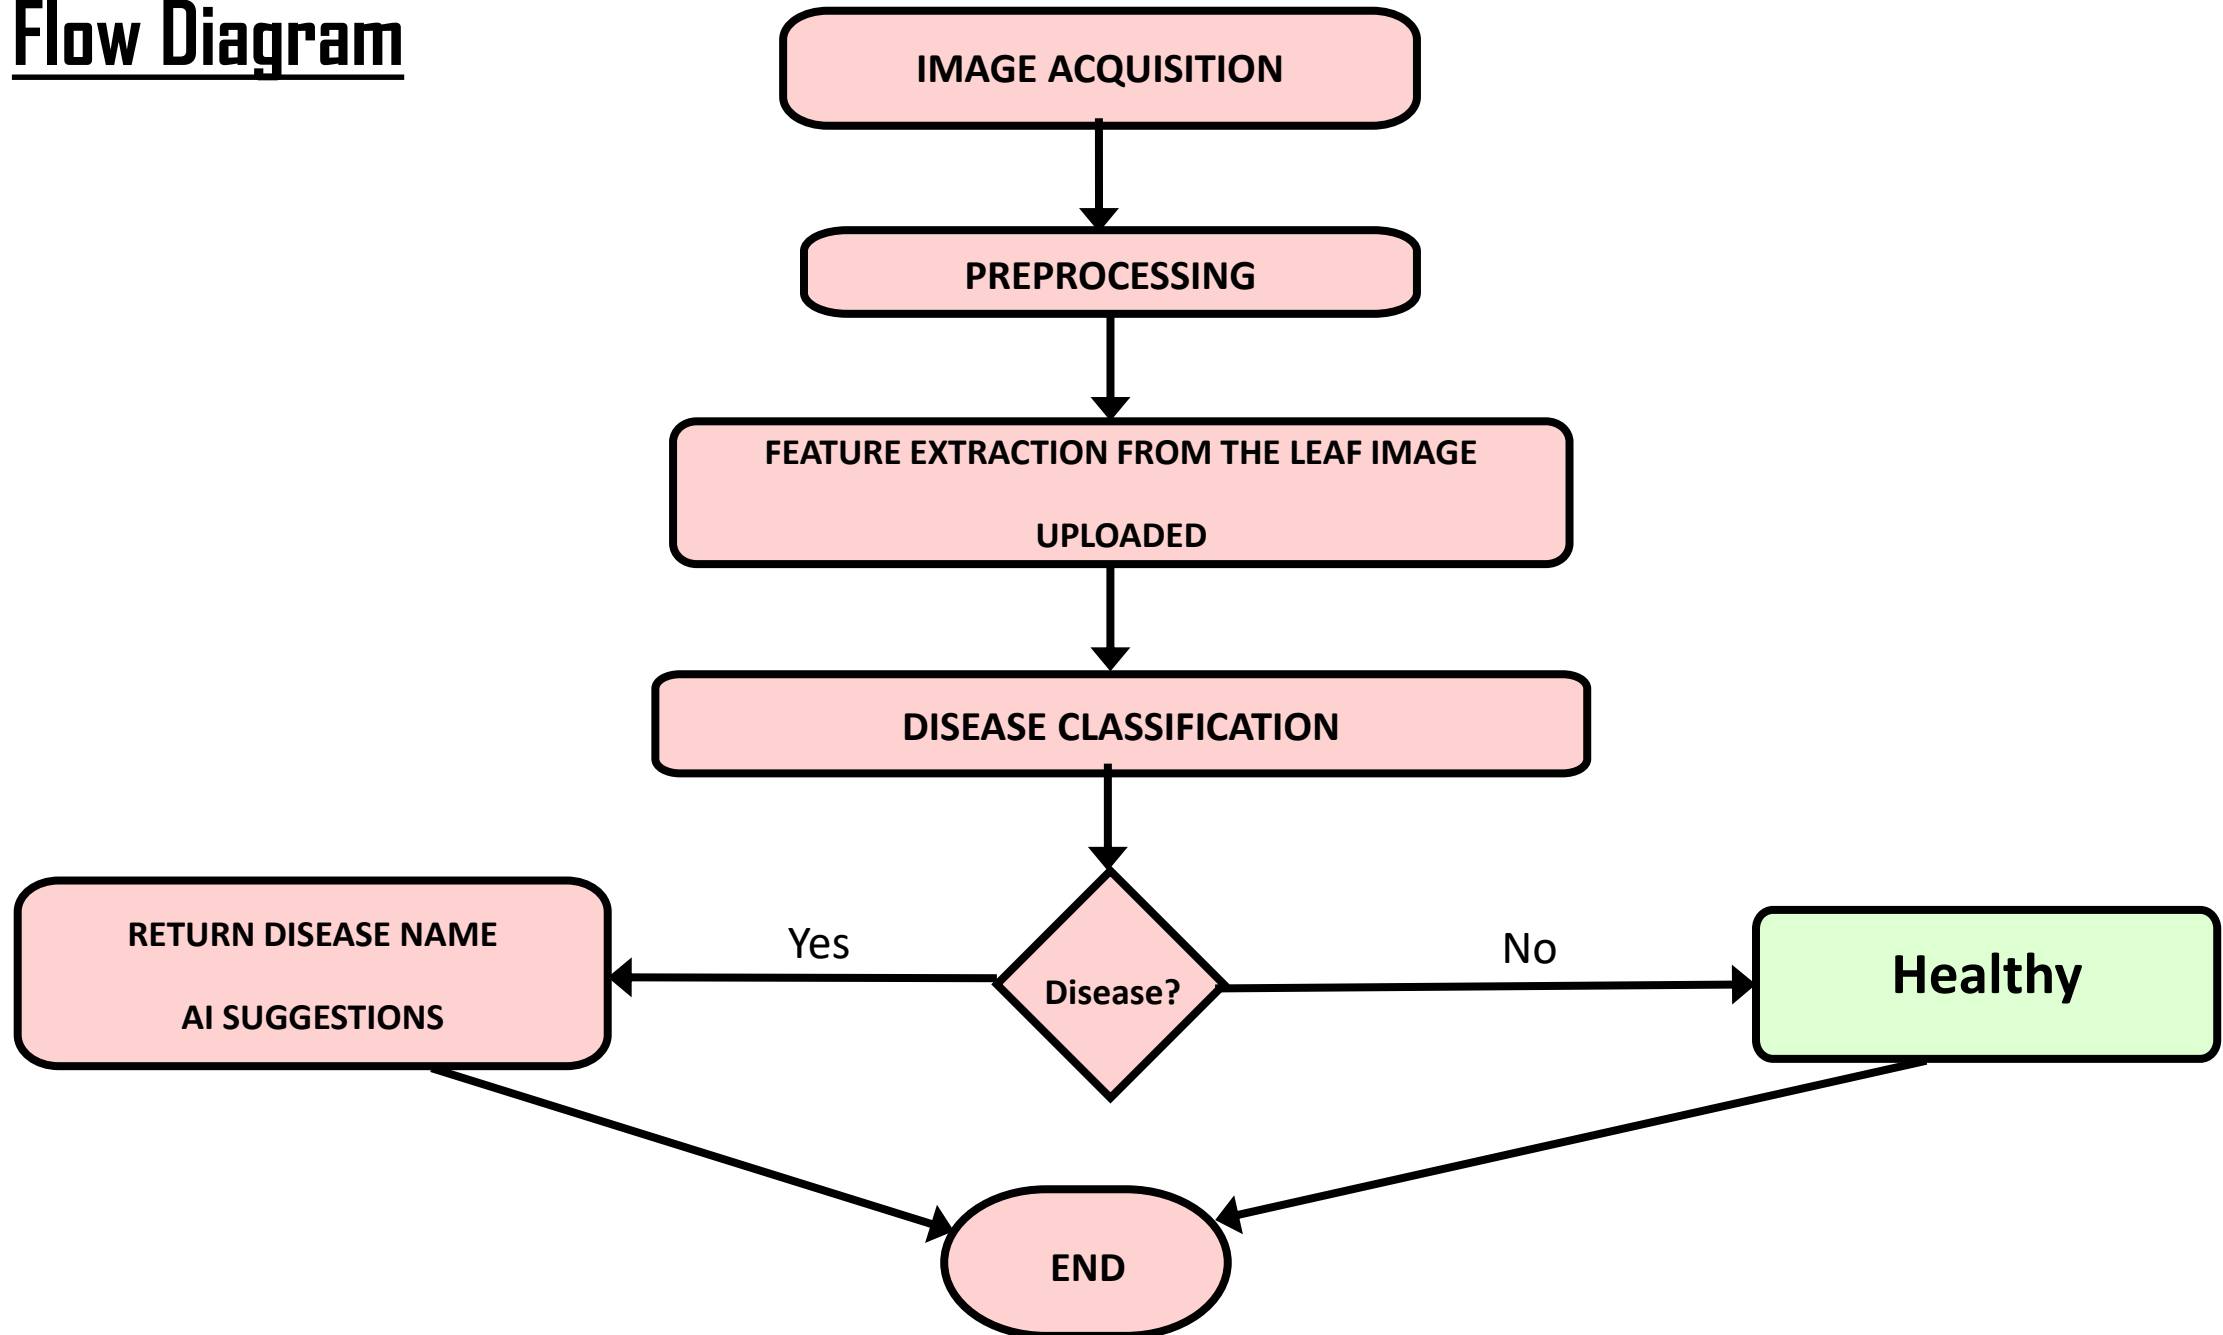
Save Class Indices: The `class_indices` dictionary is saved as `class_indices.json` for later use.

Save the Trained Model: The trained model is saved to a file (`plant_disease_prediction_model.h5`), allowing it to be loaded for future predictions.

We'll use the file i.e. **.h5** file in our frontend so that user interact with our site and upload the images then the model predict whether the disease was there or not if the disease is detected then along with prediction an AI suggestion message automatically generated.

(we used Gemini 1.5 API (text-text) so it'll generate suggestions accordingly) . And we used Streamlit a Python Library to implement frontend and logical operations.


Flow Diagram




Outputs:


Plant Disease Detection

localhost:8501




Uploaded plant leaf

 **Disease Detected**


 **Northern Leaf Blight: Long gray-green lesions on maize leaves.**

Confidence: 78.0%

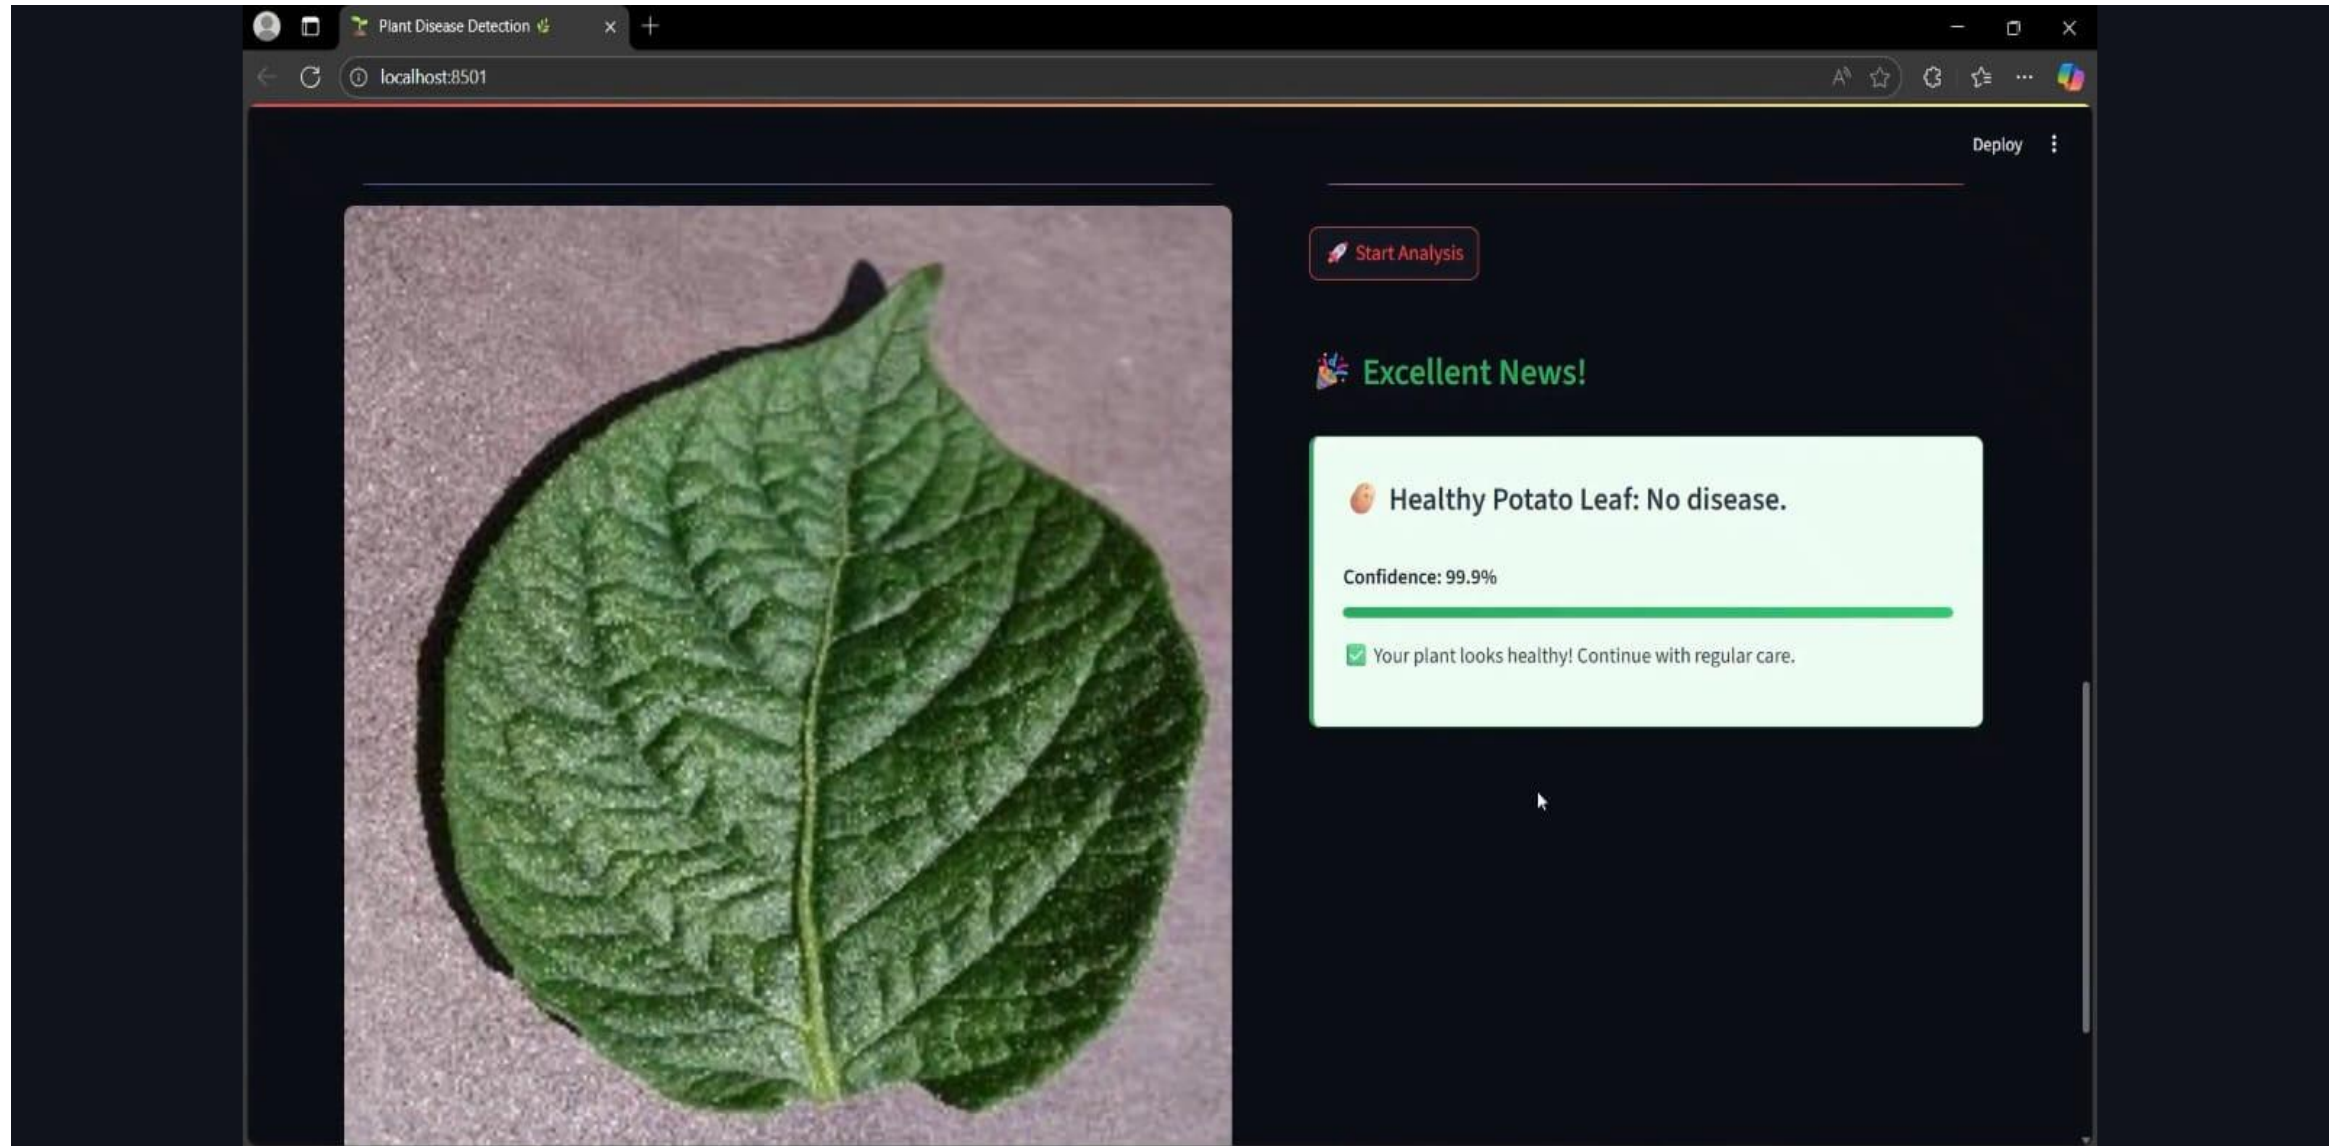
 **Treatment Recommendations**

* Use certified disease-free seeds: This is the best way to start with healthy plants and avoid the blight from the beginning.

- Rotate your crops: Don't plant corn in the same field year after year. Switch to a different crop to disrupt the blight's life cycle.
- Remove and destroy infected plant material: Get rid of any leaves or stalks showing signs of the disease immediately. Don't compost them; dispose of them properly to prevent the disease from spreading.

 **Note:** These are AI-generated suggestions. For severe cases, please consult agricultural experts.

Outputs:



The screenshot displays a web browser window with the title "Plant Disease Detection" and the address bar showing "localhost:8501". The application interface has a dark background. On the left, a large image of a healthy green potato leaf is shown. To the right of the image is a red "Start Analysis" button. Below this button, the text "Excellent News!" is displayed in green. A light green notification box contains the following information:

- 🍅 Healthy Potato Leaf: No disease.
- Confidence: 99.9%
- 🟢 Your plant looks healthy! Continue with regular care.

References:

- Real-Time Plant Leaf Disease Detection using CNN and Solutions to Cure with Android App
<https://ieeexplore.ieee.org/document/10425034>
- Using Deep Learning for Image-Based Plant Disease Detection
<https://arxiv.org/abs/1604.03169>
- PlantXViT: Explainable Vision Transformer Enabled CNN for Plant Disease Identification
<https://arxiv.org/abs/2207.07919>
- LeafGAN: An Effective Data Augmentation Method for Practical Plant Disease Diagnosis
<https://arxiv.org/abs/2002.10100>
- A hybrid framework using CNNs and Vision Transformer (ViT)
<https://link.springer.com/article/10.1007/s40747-024-01764-x>
- Soybean Disease Detection via Interpretable Hybrid CNN-GNN (CNN)
<https://arxiv.org/abs/2503.01284>

THANK YOU

~ SAI BHASKAR NANDURI